

# Backend & Infrastructure Take-Home Test

## Org Chart Service & Containerized Deployment

---

### Mission

Build a minimal backend service using **FastAPI** that manages **multiple organization charts**, including robust hierarchical operations and multi-tenant storage. The goal is to demonstrate clear backend logic, data-modeling that scales to *10,000 distinct org charts*, containerized deployment with **Docker** and **PostgreSQL**, and configuration-management best practices. Your solution should expose CRUD operations, hierarchy queries, and deletion rules that keep every chart valid at all times.

---

### Time Expectation

- **Target:**  $\leq 2$  hours for an experienced developer to finish core functionality, plus  $\frac{1}{2}$  hour for documentation & polish (total  $\leq 2.5$  h).
  - **Time Log:** Include an approximate breakdown of the time you spent on each major area (API logic, containerization, database setup, performance demo, documentation, etc.).
- 

### Contact for Questions

For clarification at any point, email [meni@thebizconquest.com](mailto:meni@thebizconquest.com).

---

# Requirements

## 1. Service Functionality (Multi-Org API with Hierarchy)

### Data Model

Table	Columns	Notes
org_charts	id (PK, int), name (string)	Represents an independent organization.
employees	id (PK, int) org_id (FK → org_charts.id) name (string), title (string), manager_id (FK → employees.id, nullable)	Self-referencing hierarchy. (manager_id=NULL means CEO).

### Core Rules

1. **Valid Hierarchy** – no cycles; every employee (except CEO) must have a manager in the same org.
2. **Delete with Re-parenting** – Deleting a non-CEO employee *must* atomically re-assign **all** of their direct (and indirect) reports to the deleted employee's own manager.
3. **CEO Protection** – The CEO **cannot be deleted**. To change the CEO, implement **PUT /orgcharts/{org\_id}/employees/{employee\_id}/promote** (or similar) that promotes an existing employee to CEO, automatically demoting/re-parenting the previous CEO.

### API Endpoints (minimum)

```
Unset
POST  /orgcharts          # create a new org chart
GET   /orgcharts          # list org charts

POST  /orgcharts/{org_id}/employees # create employee in an org
GET   /orgcharts/{org_id}/employees # list employees in org
GET   /orgcharts/{org_id}/employees/{id}
PUT   /orgcharts/{org_id}/employees/{id}
DELETE /orgcharts/{org_id}/employees/{id}
```

```
# choose ONE hierarchy view:
GET    /orgcharts/{org_id}/employees/{id}/direct_reports
      or
GET    /orgcharts/{org_id}/employees/{id}/manager_chain

# CEO replacement
PUT    /orgcharts/{org_id}/employees/{id}/promote    # makes {id}
the new CEO
```

Use clear, conventional HTTP status codes (400 for validation errors, 409 for constraint violations, etc.).

### Scale Demonstration (10 k org charts)

- Provide a **seed script** (bash/python) that creates **10,000 org charts**, each with **5-15 employees** (CEO & 1-2 layers of hierarchy is sufficient).
- After seeding, the following queries **MUST finish in <2 seconds** on a typical developer laptop:
  - `GET /orgcharts/{org_id}/employees` (random org)
  - your chosen hierarchy endpoint (`direct_reports` or `manager_chain`)

A simple timing log in your README is fine, no formal benchmarks required.

## 2. Containerization & Deployment

- **Dockerfile** – builds your FastAPI app.
- **docker-compose.yml** – launches:
  - `api` service (FastAPI, uvicorn)
  - `db` service (PostgreSQL ≥14)
- Use `ENV` variables for DB host, port, user, pwd, DB name.
- Include an **init script** or **alembic migration** that runs automatically on container start-up.

## 3. Configuration for Environments

Show how *dev* vs. *production* settings would differ (e.g., `DEBUG`, DB creds) via environment variables. Mention how you would store secrets (Docker secrets, vault, AWS SM) in production.

## 4. Basic Security & Best Practices

- **No authentication** needed for the test, but structure code so it could be added later.
- Validate foreign keys and prevent cycles on **POST/PUT**.
- Never commit secrets.

## 5. Documentation & Explanation

Create a concise **README.md** covering:

1. **Quick-start** with docker-compose (build & run).
2. **DB initialization** & seeding the 10 k org charts.
3. **Hierarchy logic** – which endpoint you implemented and why.
4. **Deletion & CEO rules** – describe the transaction.
5. **Scaling notes** – indices you added, why queries remain fast.
6. **Performance evidence** – paste short timing output from your laptop.
7. **Potential enhancements** (as applicable).
8. **AI usage** – one line on where it helped, if applicable.
9. **Open issues** – add a list of missing parts, known bugs, etc.

---

## Deliverables

- **Source Code** – FastAPI app & models.
- **Container Artifacts** – **Dockerfile**, **docker-compose.yml**, migration/DDL scripts.
- **Seed Script** – to populate 10 000 org charts.
- **Documentation** – **README.md** & time log.
- **Short video demo** – short recording of you explaining how your solution successfully performs at scale.
- **Bonus points** – **only** if you are confident your solution works, feel free to add extra things (UI/authentication/data analysis query) for extra credit.

**Where to submit:** push to a GitHub repo and send the link to [meni@thebizconquest.com](mailto:meni@thebizconquest.com)

---

## Evaluation Criteria

Area	What We Look For
Backend Logic	Correct hierarchy operations, CEO protection, cycle prevention.
Multi-Tenant Design	Clean org isolation, efficient queries, indices.
Performance Demo	Queries on 10 k org charts meet target times.
Dockerization	Containers build & start with one command; env-var config.
Code Quality	Structure, readability, naming, typing, comments.
Documentation	Clear, minimal steps from clone → running tests.
Practicality	Completes within stated time, avoids unnecessary complexity.
Explanations	You must be able to explain all your decisions (AI said so is not a good explanation)

---

## Suggested Road-map

- Ask your favorite AI developer tool to help create a structure for you
- Make sure you are happy with suggestions

Python

```
while not completed:  
    implement missing parts/fix bugs/add your logic/consult AI
```

## Tips

- Finish **core rules first**, then aim for clean code.
- Use migrations to avoid "it works on my machine" schema drift.
- Keep containers stateless; rely on DB for persistence.
- AI tools are **strongly encouraged**, use them smartly to stay within time.

**Good luck—we're excited to see what you build!**